

# CS4110-High-Performance Computing (HPC) — CCP Deliverable 1 Report



Performance Analysis and Baseline Profiling of  
KLT Feature Tracker

## Complex Computing Problem:

KLT Feature Tracker Acceleration on GPUs

## Submitted by:

Muhammad Nasir Bilal

23I-0659

Ahmed Asim

23I-0070

Github: <https://github.com/Nasir-Bilal/KLT-GPU-Accelerato>

## Instructor:

Dr. Imran Ashraf

Department of Computer Science

FAST National University of Computer and Emerging Sciences  
(NUCES) Islamabad, Pakistan

# 1. Introduction

The KLT (Kanade-Lucas-Tomasi) feature tracker is a fundamental computer vision algorithm used to detect and track distinctive points (features) across video frames. This project focuses on accelerating the KLT feature tracker using GPUs, with the goal of improving computational performance while maintaining correctness and accuracy.

Key points:

- Feature tracking detects corners and edges to follow motion in videos.
- KLT uses optical flow assumptions and a multi-scale pyramid approach for robustness.
- GPU acceleration aims to parallelize compute-intensive parts of the algorithm to exploit data-level parallelism.

## Source Code Explanation

The V1 source code implements the sequential KLT feature tracker. The main workflow involves loading input frames, detecting strong corner features, tracking these features across frames using optical flow, and refining results with optional pyramid scaling.

Utility functions handle image I/O, memory management, and storing feature lists/tables. Profiling wrappers measure execution time to identify computational hotspots for GPU acceleration.

**Examples:**

- **Example 1:** Finds the 100 best features in a single image and tracks them to the next frame. Saves feature locations to text and PPM files, and prints them to the screen.
- **Example 2:** Similar to Example 1, but replaces any lost features in the second image to maintain a constant number of tracked features.
- **Example 3:** Tracks 150 features across multiple frames (up to 10). Stores results in a feature table and writes each frame's feature list to PPM files. Sequential mode speeds up processing.

- ## 2. Profiling and Hotspot Identification

1. **gprof** identifies CPU hotspots and function-level bottlenecks.
2. **KCacheGrind** helps visualize complex call relationships and pinpoint parallelizable code.

[illegible]

```
flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative    self           self             total
time  seconds      seconds   calls   ms/call  ms/call  name
45.45    0.05      0.05        63     0.79    0.79  _convolveImageHoriz
27.27    0.08      0.03   2069270    0.00    0.00  _interpolate
 9.09    0.09      0.01    6235     0.00    0.00  _computeGradientSum
 9.09    0.10      0.01     63     0.16    0.16  _convolveImageVert
 9.09    0.11      0.01        1    10.00   12.86  _KLTSelectGoodFeatures
```

- The function `_convolveImageHoriz` alone accounts for approximately 45% of the total execution time, making it the primary hotspot of the application.
- The `_interpolate` function follows as the second heaviest contributor, consuming around 27% of runtime.

- While the remaining functions are less significant individually, they still contribute non-trivially to execution time.

**Insight:** Focusing optimization efforts on `_convolveImageHoriz` and `_interpolate` has the potential to yield the most substantial performance gains, potentially improving overall runtime by up to ~70%.

## 2. Calls vs time

- `_interpolate` is called 2,069,270 times. While the measured time per call is small, the aggregate cost is significant due to the call volume; therefore, optimizing the contexts that repeatedly invoke `_interpolate` may produce larger overall speedups than micro-optimizing `_interpolate()` alone.
- `_KLTSelectGoodFeatures` is called once, total time 0.01 s → not frequent, so impact is limited, but minor improvements are applied.

## 3. Where GPU acceleration might help

Convolution functions (`_convolveImageHoriz`, `_convolveImageVert`) are data-parallel and good candidates for GPU offload. `_computeGradientSum` and `_computeIntensityDifferenceLightingInsensitive` also contain large numbers of similar windowed operations and can be parallelized on CPU (SIMD/OpenMP) or offloaded to GPU. Although `_interpolate()` is inexpensive per call, its very high call count means caller-level optimizations (caching precomputed interpolations, inlining, loop vectorization, parallelizing outer loops) can be more effective; GPU acceleration remains attractive, especially when kernel designs avoid redundant interpolations (or use hardware/texture interpolation).

## 4. Takeaways for optimization

Prioritize `_convolveImageHoriz` and `_convolveImageVert` for the largest gains. Also optimize `_KLTSelectGoodFeatures` and caller functions such as `_computeGradientSum` — e.g., cache windowed interpolations, inline `_interpolate()`, and apply SIMD/OpenMP — to reduce the cost of repeated `_interpolate` calls. Use GPU offload or SIMD where appropriate, and re-profile after each change to measure speedup and find new hotspots..

# 3. Amdahl's-law

Formula:

Formula (single component):

$$\text{Speedup} = 1 / ((1 - P) + P / S)$$

where P = fraction of work sped up, S = speedup factor for that part.

General form (multiple components):

$$\text{Overall Speedup} = 1 / \sum_i (f_i / s_i)$$

where  $f_i$  = fraction of total time spent in component i ( $\sum f_i = 1$ ), and  $s_i$  = speedup applied to component i (use  $s_i = 1$  if unchanged).

### Baseline fractions (from gprof):

- `_convolveImageHoriz`: 45.45%
- `_interpolate`: 27.27%
- `_computeGradientSum`: 9.09%
- `_convolveImageVert`: 9.09%
- `_KLTSelectGoodFeatures`: 9.09%

### Key scenarios (theoretical overall speedups):

- *ConvolveHoriz* ×10 → ~**2.11×** overall
- *Cache/interpolate effective* ×5 → ~**1.72×** overall
- *ConvolveHoriz* ×10 + *Interpolate* ×5 → ~**3.42×** overall

These are ideal upper bounds assuming perfect scaling and no extra overhead; always re-profile after implementing changes and confirm memory/accuracy tradeoffs.

